Threat-Adaptive Interdiction of Taiwan's Crude Oil Supply

CDT Dennis Pezan
Advisor: LTC Matthew Webb
Department of Mathematical Sciences, United States Military Academy, West Point, NY

Project page

INTRODUCTION

Taiwan imports nearly all of its crude oil by sea, through chokepoints far from home. If China wanted to squeeze Taiwan short of invasion, which disruptions would hurt most?

Applied Problem

- About 794,000 barrels of crude a day, nearly all by sea (EIA, 2025).
- 61% came through the Strait of Hormuz in 2025.
- All three crude ports are on the west coast, facing the mainland.
- About 140 days of crude in storage (Ministry of Economic Affairs, 2026).

Research Questions

- RQ1. With a fixed effort budget, which disruptions reduce Taiwan's crude supply the most?
- RQ2. Can the Threat-Adaptive Network Interdiction Model (TANIM; Webb and Caplice, 2025), built on Ukrainian rail, be applied to sea lanes?

THE NETWORK

The Taiwan Crude Oil Supply Network consists of 61 nodes and 83 edges, from four sources (Gulf, U.S., Oman, West Africa) to Taiwan's three refineries.

- Nodes represent sources, loading terminals, straits and canals, ocean junctions, Taiwan approaches and ports.
- Edges represent supply links, sea routes and port entries. Each edge has a capacity (% of Taiwan's demand it can carry), a cost for China to interdict it, and a chance p that the interdiction succeeds.

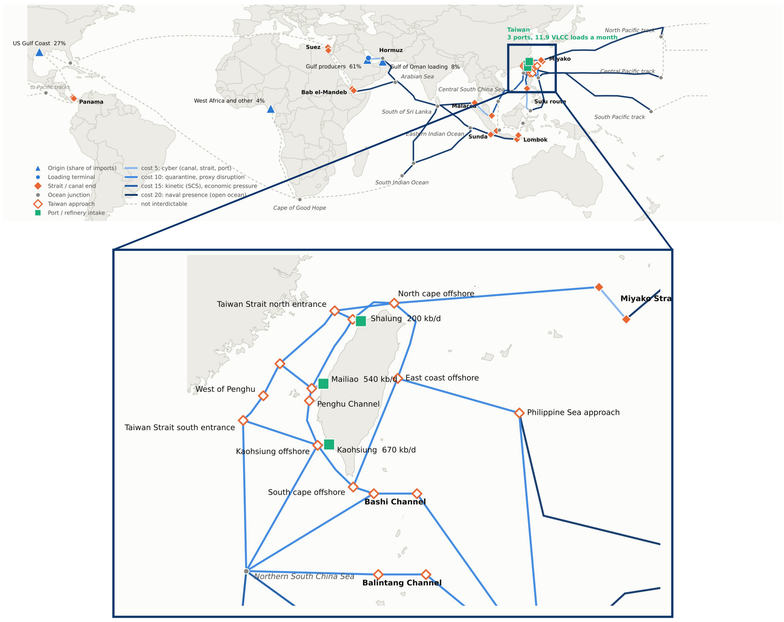


Figure 1. The Taiwan Crude Oil Supply Network (61 nodes, 83 edges). Inset: Taiwan's approaches; all three crude ports face the mainland.

METHODOLOGY

Game Setup

In a Stackelberg (leader-follower) game, one player commits to a move first and the other responds knowing it. China leads by choosing which edges to interdict; Taiwan follows by rerouting its crude to deliver as much as it can (exact max-flow).

- China's budget is fixed at 100. Each interdictable edge has one action with its own cost (5 to 20) and success chance p.
- Score: Taiwan's expected crude delivered, as % of demand. Lower is better for China.

How the RL Agent Plays

- It sees every edge's capacity, cost and p, and the budget left.
- It picks one edge at a time; edges it cannot afford or that carry no flow are masked.
- When the budget is spent, each action succeeds with its p and Taiwan reroutes.

Model and Evaluation

TANIM trains an RL attacker (graph neural network and pointer network, Maskable PPO) on a high-performance computing cluster. It is compared with:

Heuristic	What it is
Random play	Random edges until the budget is spent: the floor
Greedy	Biggest cut per unit of effort, one edge at a time, then swaps
Min-cut	Spends the budget where the network is thinnest
Local search	Simulated annealing over plans within the budget
Optimal (SAA)	Mixed-integer program on sampled outcomes; within 2 points
RL agent	Being tested: almost instant, but how close to optimal?

Table 1. Benchmarks for the RL agent

RESULTS

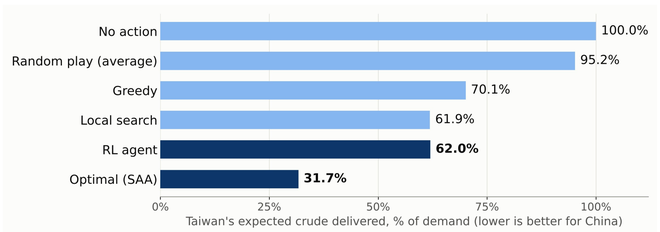


Figure 3. Base case: Taiwan's crude after China spends 100

Each scenario changes one input; Taiwan keeps (% of demand) after China spends 100:

Scenario	Greedy	Optimal	RL agent
Base case	70.1%	31.7%	62.0%
Costs 50% lower	50.4%	12.9%	*Pending*
Costs 50% higher	74.4%	52.8%	*Pending*
Taiwan needs 20% more	65.2%	31.7%	*Pending*
Taiwan needs 20% less	76.7%	31.7%	*Pending*
Hormuz fully shut	49.3%	27.2%	*Pending*

Table 2. Overall results (lower is better for China)

WHAT WE FOUND

We trained the TANIM attacker on Taiwan's network and scored every saved checkpoint on the real numbers. It settled on one plan: seal the entrances to the South China Sea, pressure the two largest suppliers, and hit the two largest ports. Taiwan keeps 62% of its crude, against 70% under greedy. The exact optimum is very different: a ring of ten quarantines around Taiwan's approaches that leaves Taiwan about 32%.

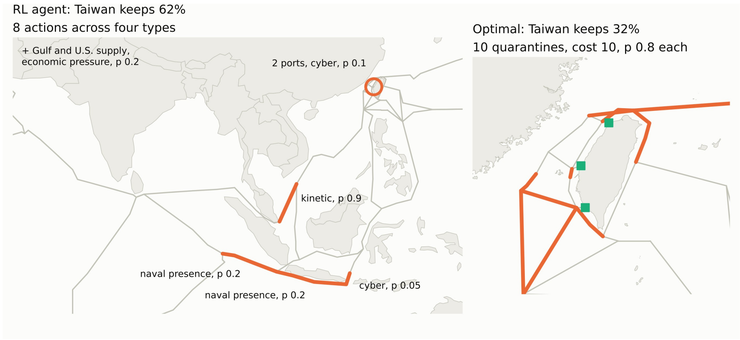


Figure 2. China's plan from the RL agent (left) and the optimal plan (right)

What Was Unexpected

- The best plan is a ring. Each quarantine alone cuts nothing, because Taiwan reroutes, so greedy, local search and the RL agent all missed it.
- It is close to Taiwan and high-odds (p 0.8), not a spread of long shots far away.
- The agent locked onto its 62% plan within its first 100,000 of 5 million training steps.

INSIGHTS

How Might This Generalize?

- Any network where an attacker has a fixed budget and uncertain odds: rail, power, supply chains. Only the edge list changes.

What New Questions Do We Have?

- Can the RL agent learn a ring, where no single step pays off?
- How much do the success chances p drive the answer?
- Which edges should Taiwan protect first to break the ring?

References

1. U.S. EIA, Country Analysis Brief: Taiwan, 2026.
2. Webb and Caplice, Solving maximum flow network interdiction using reinforcement learning, SSRN 5392397, 2025.
3. Cormican, Morton and Wood, Stochastic network interdiction, Operations Research, 1998.

ACKNOWLEDGEMENTS

Thank you to LTC Matthew Webb for advising this project and for TANIM, which it builds on, and to CDT Van Kley for collaborating on the shared code.

Keep Clear for Marking